

Anirudh Hegde

Data Scientist | Machine Learning Engineer

anirudh.hegde07@gmail.com | 9999982364 | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

SUMMARY

Data Scientist and ML Engineer with 3+ years of experience across automotive, SaaS, and manufacturing. Built end-to-end ML systems including churn prediction (XGBoost, 0.71 AUC-ROC), bilingual NLP classifiers (Vertex AI), and A/B testing frameworks. MS in Data Analytics Engineering, Northeastern University (GPA 3.92). Based in Gurugram, open to relocation across India.

EXPERIENCE

Data Analyst | Ipserlab LLC, USA

Jul 2025 – May 2026

- Analysed and documented end-to-end RFQ workflows, mapping buyer and manufacturer data requirements into structured specifications that guided product and backend development decisions
- Managed and maintained the PostgreSQL database, implementing schema changes based on evolving product requirements and ensuring data structures aligned with platform functionality
- Designed and generated synthetic datasets for backend testing, enabling the engineering team to validate data pipelines, test edge cases, and ensure system reliability before production deployment
- Collaborated cross-functionally with the backend team to define data requirements across core platform features, translating business needs into actionable data models and supporting robust pipeline development

Assistant Manager, Powertrain Planning | Maruti Suzuki India Limited, India

Jun 2022 – Jul 2023

- Conducted competitive analysis across 8 powertrain variants and 240+ models, synthesizing technology trends and regulatory shifts into data-backed roadmap recommendations presented to and adopted by senior leadership
- Owned end-to-end CAFE compliance analytics for the full vehicle fleet; modeled fuel efficiency projections across engine variants, identified optimization opportunities, and contributed to a 10% improvement in projected fuel efficiency targets
- Developed engine-wise sales and carbon emission forecasts for the full Maruti fleet through 2027, and built a live Power BI dashboard tracking weekly actuals against projections; gave executive leadership a single view to drive CAFE compliance strategy
- Represented Maruti Suzuki at Auto Expo 2023 and India Energy Week (G20 Summit India); led internal workshops on analytical tools and methodologies for the R&D department

Data Analyst | Godrej & Boyce Mfg Co Ltd, India

Jan 2021 – Jun 2022

- Analysed project performance data across 3-4 concurrent government B2B contracts, tracking cost variance, timeline adherence, and resource utilization to identify risks early and maintain delivery within budget and schedule
- Built a weekly Power BI dashboard consolidating completion status, progress metrics, and site issues across all active projects; replaced manual reporting and gave senior management a real-time view for faster intervention
- Conducted sales pipeline analysis and managed client data across potential accounts; identified high-value opportunities and supported revenue generation efforts through data-driven client prioritization
- Tracked and reported project financials weekly using Excel, analysing budget vs actuals across procurement, labour, and operations; findings presented in client meetings to drive alignment and maintain satisfaction across all government accounts

PROJECTS

AI SaaS Retention & Competitive Intelligence Platform | *dbt Core, PostgreSQL, Python, XGBoost, Streamlit, GitHub Actions* [Link](#)

- Built an end-to-end analytics pipeline ingesting 500K+ Yelp reviews and Product Hunt API data into a layered PostgreSQL warehouse (raw to staging to marts) using dbt Core; uncovered 70% month-1 churn and an 8x survival gap between Power Users and New Users
- Trained an XGBoost churn prediction model achieving 0.71 AUC-ROC; implemented 42 automated dbt tests with GitHub Actions CI and built a Streamlit dashboard visualizing cohort retention and survival curves

E-commerce Conversion Optimization | *Python, SciPy, Pandas, Plotly, Streamlit* [Link](#)

- Designed and executed an A/B test for 2,000 users using Chi-Square and Fisher's Exact tests to evaluate conversion impact of a new product page design across region and device segments
- Identified a critical segment interaction effect through Fisher's Exact Test where the new design dropped conversion from 18.4% to 0% for Europe + Tablet users ($p=0.0151$); presented findings with confidence intervals and recommended targeted rollout over full deployment, preventing a revenue-damaging global launch

Automated Bilingual Complaint Classification System | *Python, Airflow, Docker, GCP, Vertex AI, BigQuery, mBERT, GitHub Actions* [Link](#)

- Built end-to-end MLOps pipeline for bilingual complaint classification; designed Airflow DAGs for preprocessing, PII anonymization, and schema validation on 500K+ CFPB complaints, with data drift detection using cosine similarity on sentence embeddings
- Fine-tuned mBERT and XLM-RoBERTa models achieving F1 scores of 0.69 and 0.68; deployed to Vertex AI with automated CI/CD pipelines and Fairlearn bias detection with real-time Slack monitoring

SKILLS

Machine Learning & AI:	Python, Scikit-learn, XGBoost, TensorFlow, Random Forest, Logistic Regression, K-Means Clustering, NLP
Data Science & Analytics:	Survival Analysis, A/B Testing, Hypothesis Testing, Churn Modeling, RFM Analysis, Feature Engineering
MLOps & Data Engineering:	Airflow, Docker, GCP (Vertex AI, BigQuery, Cloud Run), GitHub Actions, dbt Core, PostgreSQL, SQL
Visualization & BI:	Streamlit, Tableau, Power BI, Plotly, Matplotlib, Seaborn

EDUCATION

Northeastern University, Boston, Massachusetts

May 2025

Master of Science in Data Analytics Engineering, GPA: 3.92

Guru Gobind Singh Indraprastha University, New Delhi, India

Sept 2020

Bachelor of Technology in Mechanical and Automation